

Exercise set #5 (16 pts)

- The deadline for handing in your solutions is October 25th 2022 20:00.
- Return your solutions (one .pdf file and one .zip file containing Python code) in MyCourses (Assignments tab). Additionally, submit your pdf file also to the Turnitin plagiarism checker in MyCourses.
- Check also the course practicalities page in MyCourses for more details on writing your report.

1. PageRank for directed networks (16 pts)

PageRank, a generalization of eigenvector centrality for directed networks, is used by *e.g.* Google to determine the centrality of web pages. If we consider a random walker that with probability d moves to one of the successors of the current node (i.e. nodes linked *to* by the current node) and with probability $1 - d$ teleports to a random node, PageRank of each node equals the fraction of time the random walker spent in that node. In this exercise, we investigate the behavior of PageRank in a simple directed model network (see fig. 1). To get started, you can use the provided Jupyter notebook.




Figure 1: A simple directed network.

- a) (1 pt) **Load the network** given in file `pagerank_network.edg` and, as a sanity check, **visualize** it with `nx.draw`.

Hint: To load the directed network, use parameter `create_using=nx.DiGraph()` when reading the edge list. NetworkX visualization of directed graphs is somewhat ugly but sufficient for the present purposes. In fact, the spring layout algorithm in NetworkX,

which is its default algorithm for computing node positions, works only well with undirected graphs, so for computing the layout, it's better to feed the algorithm the undirected version of the network. In addition, the algorithm can give different results on different runs, so it may be useful to plot the network a few times until the result looks good.

b) (4 pts) **Write a function that computes the PageRank** on a network by simulating a random walker. In more detail,

1. Initialize the PageRank of all nodes to 0.
2. Pick the current node (the starting point of the random walker) at random.
3. Increase the PageRank of the current node with 1.
4. Select the node, to which the random walker will move next:
 - * Draw a random number $p \in [0, 1]$.
 - * If $p < d$, the next node is one of the successors of the current one. Pick it randomly.
 - * Else, the random walker will teleport. Pick the next node randomly from all the network nodes.
5. Repeat steps 3-4 N_{steps} times.
6. Normalize the PageRank values by N_{steps} .

Use your function to compute PageRank in the example network. **Visualize** the result on the network: update your visualization from a) by using the PageRank values as node color values. **Compare** your results with `nx.pagerank` by plotting both results as a function of node index.

Hints: The damping factor is normally set to $d = 0.85$. $N_{steps} = 10000$ is a reasonable choice.

c) (4 pts) The above algorithm is a naive way of computing PageRank. The actual algorithm behind the success of Google, introduced by its founders, Larry Page and Sergey Brin, is based on power iteration [1]. The power iteration can be shown to find the leading eigenvector for the “Google matrix” (or other matrices) very fast under certain conditions. An intuitive way of thinking about the power iteration algorithm is to think that at time $t - 1$ you have a vector $x(t - 1)$ where each element gives the probability of finding the walker. You use the rules of the random walk/teleportation process to find out what are the probabilities of finding the random walkers at each node at time t . That is you increase the time t and calculate $x(t)$ based on $x(t - 1)$ until the vector x doesn't change any more. **Write a function that computes the PageRank** by using power iteration. In more detail,

1. Initialize the PageRank of all nodes to $\frac{1}{n}$, where n is the number of nodes in the network. That is, at the iteration $t = 0$ your PageRank vector contains the same value for each node, and it is equally likely to find the walker in each node. (Any other initialization strategy is possible as long as the sum of all elements is one, and the closer the initial vector is to the final vector the faster you will find the final PageRank values.)
2. Increase the iteration number t by one and create a new empty PageRank vector $x(t)$.
3. Fill in each element of the new vector PageRank vector $x(t)$ using the old PageRank vector $x(t - 1)$ and the formula: $x_i(t) = (1 - d)\frac{1}{n} + d \sum_{j \in \nu_i} \frac{x_j(t-1)}{k_j^{\text{out}}}$, where ν_i is the set of nodes that have a directed link ending at i , and for each such node $j \in \nu_i$, k_j^{out}

is j 's out-degree. In summary, for each node i you need to calculate their entry in the new PageRank vector $x(t)$ as a sum of two parts:

- * probability that the walker will teleport into the node $(1 - d)\frac{1}{n}$ and
- * probability that the walker will move from a neighbor j to node i . Iterate over each in-neighbor j of the node i (i.e., there is a link from j to i) and add the neighbors contribution $d\frac{x_j(t-1)}{k_j^{\text{out}}}$ to the entry of the node i in the new PageRank vector $x(t)$.

4. Repeat steps 2 and 3 $N_{\text{iterations}}$ times.

Use your function to compute PageRank in the example network and **visualize the result** on the network as in b).

Hints:

- The damping factor is normally set to $d = 0.85$.
- You can monitor the progress of the power iteration by printing out the change in the PageRank vector $\Delta(t) = \sum_i |x_i(t) - x_i(t-1)|$ after each iteration step. The change $\Delta(t)$ should be decreasing function of t . $N_{\text{iterations}} = 10$ should be more than enough in most cases.
- You can list the incoming edges to node i with the function `net.in_edges(i)`, where `net` is the network object. Alternatively, you can use the function `net.predecessors(i)`, which returns an iterator over predecessor nodes of node i .
- The sum of all elements in the PageRank vector should always equal to one. There might be slight deviations from this due to numerical errors, but much larger or smaller values is an indication that something is wrong with the code.

d) (2 pts) The Google search engine indexes billions of websites and the algorithm for calculating the PageRank needs to be extremely fast. In the original paper about PageRank [1], by Google founders Larry Page and Sergey Brin, they claim that their “iterative algorithm” is able to calculate the PageRank for 26 million webpages in a few hours using a normal desktop computer (in 1998). **Come up with a rough estimate** of how long it would take for your power iteration algorithm (part c) and naive random walker algorithm (part b) to do the same. You can assume that the average degree of the 26 million node network is small and that the power iteration converges in the same number of steps as it does for your smaller networks. For the random walk you can assume that you need to run enough steps that the walker visits each node on average 1000 times. You can also omit any considerations of fitting the large network in memory or the time it takes to read it from the disk etc. With these assumption you can simply calculate the time it takes to run the algorithm in a reasonable size network and multiply the result by the factor that the 26 million node network is bigger than your reasonable sized network.

Report all calculations and parameters you use (such as size of the network, number of steps); simply reporting the result without telling how you obtained it will not get you any points.

Hints:

- There are several ways of timing your code. You can use for example IPython's `%timeit` command or the Python `timeit` module.

- The small example network is probably going to be too small to test out the speed of your function. You should aim for a network for which it takes several seconds to run the PageRank function (note that a network size that works for the power iteration algorithm may be far too large for the random algorithm to finish in a reasonable time). You might find it useful to use network model in networkx to run your code. For example,

```
net=nx.directed_configuration_model(10**4*[5],10**4*[5],create_using=nx.DiGraph())
```

will produce network with 10000 nodes where each node has in and out degrees of 5 using the configuration model.
 - Don't feel bad if you cannot beat Larry and Sergey in speed when using Networkx and Python. These tools are not meant for speed of computation and even modern computers might not be enough to help. Also, your competition invented Google.
- e) The aim of this task is to understand how the structure of a network relates to PageRank. Remember to **justify** your answers.
- 1) What is the connection between a node's in-degree k_{in} and its PageRank (if any)? If you know only the in-degrees of two nodes (in the same network), can you tell which one has a higher PageRank?(1 p)
 - 2) What is the connection between a node's out-degree k_{out} and its PageRank (if any)? (0.5 p)
 - 3) How does PageRank change if the node belongs to a strongly connected component? (0.5 p)
 - 4) How could the information about the network's structure be used in improving the power iteration algorithm given in part c)? (0.5 p)
Hint: Are there ways to incorporate this information to make the algorithm converge faster?
 - 5) Explain why the PageRanks of nodes 3 and 4 in the plotted network are higher than that of node 5. (0.5 p)
- f) (2 pts) **Investigate** the effect of the damping factor d on the PageRank values of the network used in parts a)-c). Repeat the PageRank calculation with *e.g.* 5 different values of $d \in [0, 1]$ and **plot** the PageRank as a function of node index (plots of all values of d in the same figure). **Interpret the results.** How and why does the change of d affect the rank of the nodes and the absolute PageRank values? **Explain** what happens, when $d=0$ and when $d=1$.

In case you do not trust your implementations in b) and c), you can use PageRank values obtained with 'nx.pagerank' in this last task.

Feedback (1 pt)

To earn one bonus point, give feedback on this exercise set and the corresponding lecture latest two days after the report's submission deadline. You can find the feedback form at the Assignments tab in MyCourses.

References

- [1] S. Brin and L. Page, “The anatomy of a large-scale hypertextual web search engine,” *Computer networks and ISDN systems*, vol. 30, no. 1-7, pp. 107–117, 1998.